



Create S3 Buckets with Terraform

GA

gazzok@gmail.com

```
validate    Check whether the configuration is valid
plan        Show changes required by the current configuration
apply       Create or update infrastructure
destroy     Destroy previously-created infrastructure

All other commands:
console     Try Terraform expressions at an interactive command prompt
fmt         Reformat your configuration in the standard style
force-unlock Release a stuck lock on the current workspace
get         Install or upgrade remote Terraform modules
graph       Generate a Graphviz graph of the steps in an operation
import      Associate existing infrastructure with a Terraform resource
login       Obtain and save credentials for a remote host
logout      Remove locally-stored credentials for a remote host
metadata    Metadata related commands
modules     Show all declared modules in a working directory
output      Show output values from your root module
providers   Show the providers required for this configuration
refresh     Update the state to match remote systems
show        Show the current state or a saved plan
state       Advanced state management
taint       Mark a resource instance as not fully functional
test        Execute integration tests for Terraform modules
untaint     Remove the 'tainted' state from a resource instance
version     Show the current Terraform version
workspace   Workspace management

Global options (use these before the subcommand, if any):
-chdir=DIR  Switch to a different working directory before executing the
            given subcommand.
-help      Show this help output, or the help for a specified subcommand.
-version   An alias for the "version" subcommand.
garyking@Garys-Mac-mini ~ % pwd
/Users/garyking
```



Introducing Today's Project!

In this project, I will demonstrate using Terraform to deploy infrastructure in AWS. The goal is to configure my AWS credentials in the terminal and create and manage S3 buckets with Terraform. Finally, I will upload files to S3 using Terraform.

Tools and concepts

Services I used were: Terraform Amazon S3 AWS CLI AWS IAM

Project reflection

The project took me approximately 1.5 hours. The most challenging part was setting up Terraform for the first time. It was most rewarding to use Terraform to launch resources in AWS using the AWS CLI. It truly showcased the power of Terraform and illustrated its usecases.

I completed the project today to improve my Terraform knowledge and learn more about Infrastructure as code (IaC). This project definitely met my goals and has given me a solid base to continue my exploring and learning of Terraform.



Introducing Terraform

Terraform is a tool that helps you build and manage your cloud infrastructure using code. Instead of setting up resources manually in the AWS console or CLI, you write a script that tells Terraform exactly what you want in your cloud infrastructure, like servers, databases, and networks. Terraform then automatically builds all of this for you, using your script as a blueprint.

Infrastructure as Code (IaC) is a way to manage your IT infrastructure. Instead of manually setting up your resources e.g. creating resources one by one in the Console, you're writing configuration files in code. Things built from code are built the same way every time. You can repair and rebuild things quickly, and other people can build identical instances of the same thing. This makes managing large-scale systems more efficient, less error-prone, and way faster than doing everything manually. Terraform is a top IaC tool because it supports multiple cloud providers, so you can set up multi-cloud infrastructures that use AWS, Azure, GCP and more. It also uses a simple configuration language.

Terraform uses configuration files to define and manage infrastructure. These files describe the desired state of your infrastructure, and Terraform figures out how to achieve that state.



```
validate    Check whether the configuration is valid
plan        Show changes required by the current configuration
apply       Create or update infrastructure
destroy     Destroy previously-created infrastructure

All other commands:
console     Try Terraform expressions at an interactive command prompt
fmt         Reformat your configuration in the standard style
force-unlock Release a stuck lock on the current workspace
get         Install or upgrade remote Terraform modules
graph       Generate a Graphviz graph of the steps in an operation
import      Associate existing infrastructure with a Terraform resource
login       Obtain and save credentials for a remote host
logout      Remove locally-stored credentials for a remote host
metadata    Metadata related commands
modules     Show all declared modules in a working directory
output      Show output values from your root module
providers   Show the providers required for this configuration
refresh     Update the state to match remote systems
show        Show the current state or a saved plan
state       Advanced state management
taint       Mark a resource instance as not fully functional
test        Execute integration tests for Terraform modules
untaint     Remove the 'tainted' state from a resource instance
version     Show the current Terraform version
workspace   Workspace management

Global options (use these before the subcommand, if any):
-chdir=DIR  Switch to a different working directory before executing the
            given subcommand.
-help       Show this help output, or the help for a specified subcommand.
-version    An alias for the "version" subcommand.
garyking@Garys-Mac-mini ~ % pwd
/Users/garyking
```



Configuration files

The configuration is structured in blocks to help organize the code better. The style of organizing code into individual blocks is called modularity. Each block handles a specific piece of your setup, like a resource or a configuration, which makes it easier to read, manage, and adjust things separately without affecting the rest of your infrastructure. These are the basic building blocks of any Terraform configuration: Provider blocks specify which plugins (like AWS, Google Cloud, or Azure) Terraform uses to deploy and manage resources in those services. Resource blocks define what resources to manage, such as networks, virtual machines, or storage buckets. Each resource is individually managed and can be customized with parameters defined by the provider.



My main.tf configuration has three blocks

The first block indicates the provider "aws": Indicates that AWS is the provider. A provider is a plugin (a software add-on) that allows Terraform to interact with a specific cloud service, API, or on-premises resource. The provider translates your configurations into API calls that manage your infrastructure. The second block provisions the resource "aws_s3_bucket" "my_bucket": This block provisions an S3 bucket. "my_bucket" is an internal alias used in your code to reference this resource. The third block manages the resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block": Manages public access policies for the S3 bucket. In our case, options like block_public_acls, ignore_public_acls, block_public_policy, and restrict_public_buckets, are all set to true to stop public access to your bucket and its contents.



```
provider "aws" {  
  1 provider "aws" {  
  2   region = "eu-west-2" # Update this to the Region closest to you  
  3 }  
  4  
  5 resource "aws_s3_bucket" "my_bucket" {  
  6   bucket = "nextwork-unique-bucket-garyking-13" # Ensure this bucket name is globally unique  
  7 }  
  8  
  9 resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block" {  
 10   bucket = aws_s3_bucket.my_bucket.id  
 11  
 12   block_public_acls       = true  
 13   ignore_public_acls     = true  
 14   block_public_policy     = true  
 15   restrict_public_buckets = true  
 16 }  
 17
```



Customizing my S3 Bucket

For my project extension, I visited the official Terraform documentation to read documentation on the aws provider, where we can find setup instructions, descriptions of each available resource, best practice tips and examples, and parameters for customization. The `aws_s3_bucket` in the Terraform documentation is all about how to use Terraform to create and manage AWS S3 buckets. It covers how to configure these buckets, set their properties, and manage interactions with other AWS services.

I chose to customise my bucket by adding tags because I wanted to be able to differentiate if this file was for a development environment or production environment.



```
1 provider "aws" {
2   region = "eu-west-2"
3 }
4
5 resource "aws_s3_bucket" "my_bucket" {
6   bucket = "nextwork-unique-bucket-garyking-13"
7
8   tags = {
9     Name       = "My bucket"
10    Environment = "Dev"
11  }
12 }
13
14 resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block" {
15   bucket = aws_s3_bucket.my_bucket.id
16
17   block_public_acls       = true
18   ignore_public_acls     = true
19   block_public_policy     = true
20   restrict_public_buckets = true
21 }
```



Terraform commands

I ran 'terraform init' to Terraform init is the first command you run for a new Terraform project. It sets up your project by: Downloading necessary plugins: It grabs the specific providers Terraform needs to manage your cloud services, like AWS. Setting up the backend: Prepares where Terraform will store its record of your setup, so it knows what your cloud setup looks like and can manage it efficiently. Preparing modules: If your setup uses reusable code components, called modules, this is where the code is downloaded from their source to replace placeholders in your code. Creating a lock file: This file keeps track of the versions of everything Terraform is using, making sure that everyone on your team is using the same setup to avoid inconsistencies.

Next, I ran 'terraform plan' to create an execution plan, showing you what changes Terraform will make to my infrastructure based on my configuration file. It shows what will be created, updated, or destroyed so we can review and confirm the configuration before any real changes are made.



```
garyking@Garys-Mac-mini nextwork_terraform % terraform init
Initializing the backend...
Initializing provider plugins...
- Finding latest version of hashicorp/aws...
- Installing hashicorp/aws v5.89.0...
- Installed hashicorp/aws v5.89.0 (signed by HashiCorp)
Terraform has created a lock file .terraform.lock.hcl to record the provider
selections it made above. Include this file in your version control repository
so that Terraform can guarantee to make the same selections by default when
you run "terraform init" in the future.

Terraform has been successfully initialized!

You may now begin working with Terraform. Try running "terraform plan" to see
any changes that are required for your infrastructure. All Terraform commands
should now work.

If you ever set or change modules or backend configuration for Terraform,
rerun this command to reinitialize your working directory. If you forget, other
commands will detect it and remind you to do so if necessary.
garyking@Garys-Mac-mini nextwork_terraform %
```



AWS CLI and Access Keys

When I tried to plan my Terraform configuration, I received an error message that says 'Planning Failed' because I did not have AWS CLI installed.

To resolve my error, first I installed AWS CLI, which is The AWS CLI (Command Line Interface). This is a powerful tool that lets you manage your AWS services from your terminal. Instead of having to use the AWS Management Console, you can now run text commands from your local machine.

I set up AWS access keys to connect the CLI to the management console as the CLI needs another way of authenticating you.

```
garrying@Garrys-Mac-mini nextwork_terraform % terraform plan
Planning failed. Terraform encountered an error while generating this plan.

Error: No valid credential sources found

  with provider["registry.terraform.io/hashicorp/aws"],
   on main.tf line 1, in provider "aws":
    1: provider "aws" {}

Please see https://registry.terraform.io/providers/hashicorp/aws
for more information about providing credentials.

Error: failed to refresh cached credentials, no EC2 IMDS role found, operation error ec2imds: GetMetadata, request canceled, context deadline exceeded

garrying@Garrys-Mac-mini nextwork_terraform % aws --version
zsh: command not found: aws
garrying@Garrys-Mac-mini nextwork_terraform % curl "https://awscli.amazonaws.com/AWSCLIV2.pkg" -o "AWSCLIV2.pkg"
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left  Speed
100 39.1M 100 39.1M    0     0  35.0M      0  0:00:01  0:00:01 --:--:-- 35.0M
Password:
installer: This package requires Rosetta 2 to be installed.
Please install Rosetta 2 and then try again.
  sudo softwareupdate --install-rosetta

installer: Error - AWS Command Line Interface can't be installed on this computer.
garrying@Garrys-Mac-mini nextwork_terraform % sudo softwareupdate --install-rosetta
I have read and agree to the terms of the software license agreement. A list of Apple SLAs may be found here: https://www.apple.com/legal/sla/
Type A and press return to agree: A
2025-09-02 18:06:11.374 softwareupdate[14363:1442232] Package Authoring Error: 072-78605: Package reference com.apple.pkg.RosettaUpdatesAuto is missing installKBBytes
Install of Rosetta 2 finished successfully
garrying@Garrys-Mac-mini nextwork_terraform % curl "https://awscli.amazonaws.com/AWSCLIV2.pkg" -o "AWSCLIV2.pkg"
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload   Total   Spent    Left  Speed
100 39.1M 100 39.1M    0     0  35.0M      0  0:00:01  0:00:01 --:--:-- 35.0M
installer: Package name is AWS Command Line Interface
installer: Installing at base path /
installer: The install was successful.
garrying@Garrys-Mac-mini nextwork_terraform %
```



Lanching the S3 Bucket

I ran 'terraform apply' to apply the changes I'd written in my Terraform configuration. It creates, modifies, or deletes resources in your infrastructure based on your code. In my case, this command creates the S3 bucket in my AWS account.

If you run terraform apply before terraform init, you would run into an error because terraform init needs to set up your project first by downloading necessary plugins and setting up the state file, which is a file Terraform uses to track the current state of your infrastructure. There actually aren't any errors if you skip terraform plan. This is because terraform apply will automatically run a plan and ask for confirmation before making any changes. But, skipping the explicit terraform plan means you might miss reviewing these changes in detail. It's best practice to run a plan to catch any potential errors or unexpected results before they affect your live infrastructure.



The screenshot shows the Amazon S3 console interface. On the left is a sidebar with navigation links for Amazon S3 features like General purpose buckets, Directory buckets, Table buckets, Access Grants, Access Points, Object Lambda Access Points, Multi-Region Access Points, Batch Operations, and IAM Access Analyzer for S3. Below these are links for Blocking Public Access settings and Storage Lens. The main content area is titled 'General purpose buckets (1)' and includes a search bar and a table of buckets. The table has columns for Name, AWS Region, IAM Access Analyzer, and Creation date. One bucket is listed: 'network-unique-bucket-garyling-13' in the 'Europe (London) eu-west-2' region. Above the table are buttons for 'Copy ARN', 'Empty', 'Delete', and 'Create Bucket'. At the top of the console, there's a banner for 'Account snapshot - updated every 24 hours' and a 'View Storage Lens dashboard' button.

Name	AWS Region	IAM Access Analyzer	Creation date
network-unique-bucket-garyling-13	Europe (London) eu-west-2	View analyzer for eu-west-2	March 2, 2025, 10:28:15 (UTC+00:00)

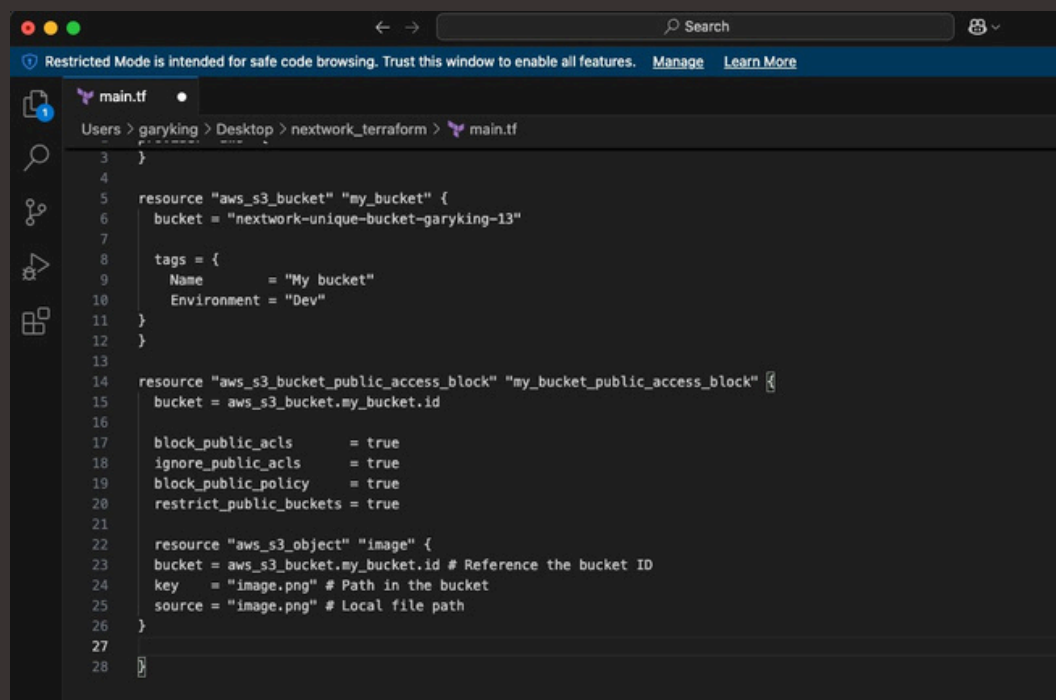


Uploading an S3 Object

I created a new resource block to add an image to my S3 bucket using Terraform.

We need to run terraform apply again because you need to run terraform apply anytime you change your Terraform configuration. This will make sure the changes are correctly applied to your infrastructure.

To validate that I've updated my configuration successfully, I logged back into my management console and into my S3 bucket. I then was able to see my image.png which I downloaded and opened to verify it worked - which it did!



```
main.tf
Users > garyking > Desktop > nextwork_terraform > main.tf
3  }
4
5  resource "aws_s3_bucket" "my_bucket" {
6    bucket = "nextwork-unique-bucket-garyking-13"
7
8    tags = {
9      Name      = "My bucket"
10     Environment = "Dev"
11   }
12 }
13
14 resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block" {
15   bucket = aws_s3_bucket.my_bucket.id
16
17   block_public_acls       = true
18   ignore_public_acls     = true
19   block_public_policy     = true
20   restrict_public_buckets = true
21
22   resource "aws_s3_object" "image" {
23     bucket = aws_s3_bucket.my_bucket.id # Reference the bucket ID
24     key    = "image.png" # Path in the bucket
25     source = "image.png" # Local file path
26   }
27 }
28 }
```